



FINAL PROJECT REPORT NEWS DATA PIPELINE

Team Members

Omar Alua - 22B030417

Sagatkyzy Firuza - 22B030425

Yesserkey Dana - 23B030349

API Selection and Justification

We use the NewsData.io High-frequency News API (<https://newsdata.io/>).

This API provides real-time news articles categorized by keywords, categories, languages, and countries, which makes the data suitable for analytical processing.

Justification

- The API fully satisfies the assignment requirements
- It is stable and well-documented
- It is frequently updated, with new news appearing every few minutes
- It provides real-world, meaningful data from trusted sources
- It returns structured data in JSON format

Kafka topic schema

In the first stage of the pipeline, a producer is implemented, which is responsible for receiving news data from an external API. NewsData.io and sending them to Kafka topic `raw_events`

Topic name: **raw_events**

message format: **JSON**

One message corresponds to one news article

```
{  
  "data": {  
    "article_id": "ea4a7165ab9c179b9833689d55b70836",  
    "link": "https://arr.news/2025/12/16/the-buloke-times-16-december-2025/",  
    "title": "The Buloke Times, 16 December 2025",  
    "description": "Out now! Buy here! I Subscribe here!",  
    "content": "ONLY AVAILABLE IN PAID PLANS",  
    "keywords": [  
      "vic",  
      "latest",  
      "buloke times",  
      "out now",  
      "news",  
      "december 2025"  
    ],  
    "creator": [  
      "The Buloke Times"  
    ],  
    "language": "english",  
    "country": [  
      "australia"  
    ]  
  }  
}
```

Message schema:

```
{  
  "article_id": "string",  
  "title": "string",  
  "description": "string | null",  
  "content": "string | null",  
  "link": "string",  
  "source_id": "string",  
  "source_name": "string",  
  "category": "string",  
  "country": "string",  
  "language": "string",  
  "pubDate": "string",  
  "ingested_at": "timestamp"  
}
```

Each message is posted to Kafka by the developer in DAG 1 (job1_news_ingestion) and later processed by the DAG cleanup team.

Cleaning Rules

Data cleaning is performed in DAG 2, which consumes raw news messages from the Kafka topic `raw_events` and prepares them for storage in SQLite

We applied the following rules:

1. Mandatory fields validation

Records must contain non-empty `article_id` and `title`.

Titles shorter than 10 characters are discarded.

2. Text normalization

Extra whitespace is removed and text fields (`title`, `description`, `content`) are normalized.

3. Handling optional fields

Optional fields (`description`, `content`) may be null and are stored as `NULL` in SQLite.

4. Category, country, and language normalization

List values are joined into comma-separated strings, default values are applied when missing, and all values are converted to lowercase.

5. Date and metadata handling

The publication date (`pubDate`) is stored as received, and a `cleaned_at` timestamp is added.

6. Duplicate prevention and error handling

Duplicate records are ignored during insertion, and malformed messages are logged and skipped without stopping the batch job.

SQLite schema (for both tables)

```
sqlite> .tables
daily_summary  news_events
```

Table: `news_events`

This table stores cleaned and normalized news articles after processing in DAG 2.

```
news_events (
    id INTEGER PRIMARY KEY AUTOINCREMENT,
    article_id TEXT UNIQUE,
    title TEXT NOT NULL,
    description TEXT,
    content TEXT,
    link TEXT,
    source_id TEXT,
    source_name TEXT,
    category TEXT,
    country TEXT,
    language TEXT,
    pubDate TEXT,
    ingested_at TEXT,
    cleaned_at TEXT
)
```

The news_events table stores cleaned and normalized news articles received from Kafka after processing in DAG 2

```
• (venv1) aluwa@ulu:/mnt/d/final_data_collection/airflow/dags$ python3 ./news_events.py
{
    "id": 1,
    "article_id": "8d4c7f0469df58f53a9885a4eb406af8",
    "title": "Thai wage growth slows amid sluggish economy",
    "description": "Salary increases in Thailand are moderating to average 4.5% across industries and businesses in 2025, slowing from the historical norm of 5%, says international consultancy Deloitte.",
    "content": "ONLY AVAILABLE IN PAID PLANS",
    "link": "https://www.bangkokpost.com/business/general/3157384/thai-wage-growth-slows-amid-sluggish-economy",
    "source_id": "bangkokpost",
    "source_name": "Bangkok Post",
    "category": "top,business",
    "country": "thailand",
    "language": "english",
    "pubDate": "2025-12-15 22:26:00",
    "ingested_at": "2025-12-16 10:57:52",
    "cleaned_at": "2025-12-16 10:57:52.448759"
}
```

The daily_summary table contains the results of daily analytics calculated in DAG 3 based on data from the news_events table. It is used to store aggregated metrics on news for the day

Table: daily_summary

This table stores aggregated daily analytics calculated in DAG 3.

```
daily_summary (
    date TEXT,
    total_articles INTEGER,
    unique_sources INTEGER,
```

```

top_category TEXT,
top_country TEXT,
average_title_length REAL
)

```

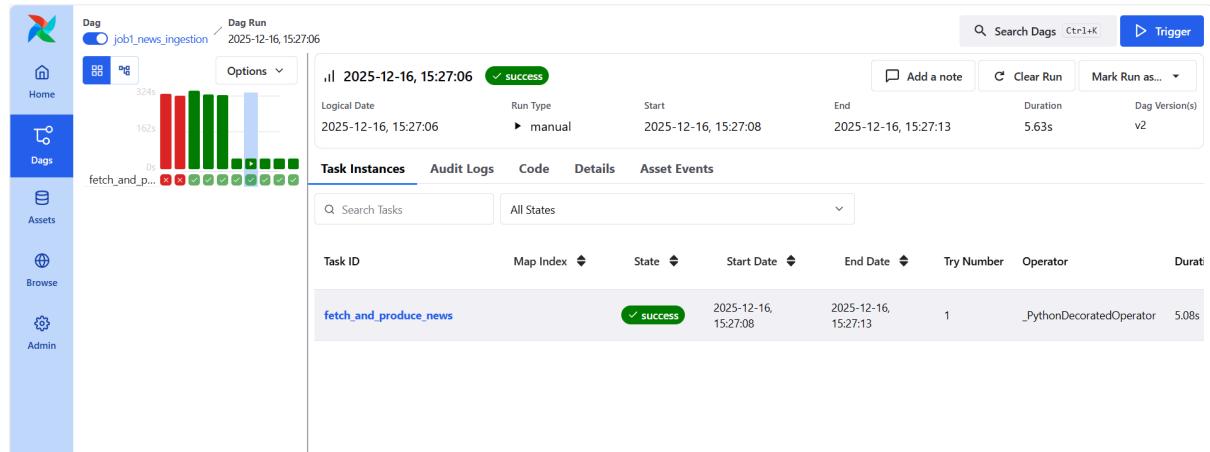
```

• (venv1) aluwa@Lulu:/mnt/d/final_data_collection/airflow/dags$ python3 ./daily_summary.py
+-----+
| Date | Total articles | Unique sources | Top Category | Top country | Average title length |
+-----+
| 2025-12-16 | 209 | 129 | top,business | united states of america | 76.22 |
+-----+

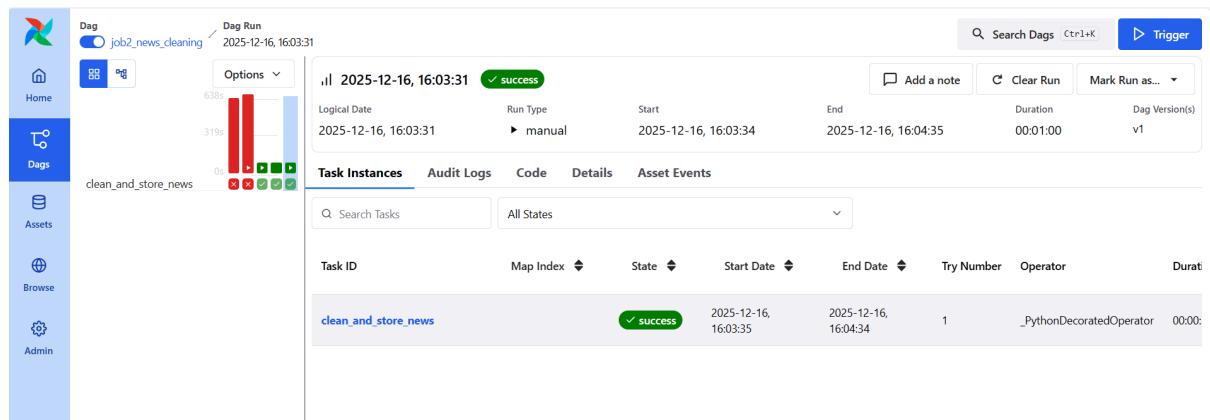
```

DAG1

Successful execution of the news ingestion DAG (job1_news_ingestion) that fetches data from the NewsData.io API and publishes raw news messages to the Kafka topic raw_events.



DAG2



Successful execution of the data cleaning DAG (job2_news_cleaning) that consumes messages from Kafka, applies cleaning rules, and stores cleaned data in SQLite.

DAG3

The screenshot shows the Airflow web interface for the DAG `job3_daily_analytics`. The top navigation bar includes a search bar for 'Search Dags' and a 'Trigger' button. On the left, a sidebar menu lists 'Home', 'Dags' (which is selected), 'Assets', 'Browse', and 'Admin'. The main content area displays the DAG run details for the logical date 2025-12-16, 15:58:17, which was run manually and completed successfully. The task instance table shows one task named `compute_daily_summary` with a state of `success`, starting at 2025-12-16, 15:58:19 and ending at 2025-12-16, 15:58:20. The duration was 0.52s and the dag version was v1.

Map Index	State	Start Date	End Date	Try Number	Operator	Duration	Dag Version
	<code>success</code>	2025-12-16, 15:58:19	2025-12-16, 15:58:19	1	<code>_PythonDecoratedOperator</code>	0.52s	v1

Successful execution of the daily analytics DAG (`job3_daily_analytics`) that computes aggregated metrics from SQLite and writes results to the `daily_summary` table.